

dog_app

May 14, 2021

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [67]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/"))
         dog_files = np.array(glob("/data/dog_images/*/"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [68]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [70]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    if len(faces) > 0:
        return True
    return False
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

the accuracy of the algorithm in human files 0.98

the accuracy of the algorithm in dog files 0.17

```
In [71]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
counter_h=0
counter_d=0

for img in human_files_short:
    humen=face_detector(img)
    if humen :
        counter_h += 1

for img in dog_files_short :
    humen=face_detector(img)
    if not humen:
```

```

        counter_d += 1

percent_in_humen_files=counter_h/len(human_files_short)

percent_in_dog_files=1-(counter_d/len(dog_files_short))

print("the accuracy of the algorithm in humen files {}".format(percent_in_humen_files))
print("the accuracy of the algorithm in dog files {}".format(percent_in_dog_files))

the accuracy of the algorithm in humen files 0.98
the accuracy of the algorithm in dog files 0.17000000000000004

```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```

In [72]: import torch
import torchvision.models as models

# define VGG16 model
VGG16 = models.vgg16(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()

print(VGG16.classifier[0].in_features)
VGG16

```

25088

```
Out[72]: VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): ReLU(inplace)
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (27): ReLU(inplace)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace)
    (2): Dropout(p=0.5)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace)
    (5): Dropout(p=0.5)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000

possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [73]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    data_transform = transforms.Compose([ transforms.Resize((224,224)),transforms.ToTensor()])
    img= Image.open(img_path)
    img_t=data_transform(img)
    batch_t = torch.unsqueeze(img_t, 0)
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    VGG16.eval()
    output= VGG16(batch_t.to(device))
    output=output.to(device)
    _, preds_tensor = torch.max(output, 1)
    preds_tensor=preds_tensor.to(torch.device("cpu"))
    preds = np.squeeze(preds_tensor.numpy())
    return preds # predicted class index
```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all

categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```
In [74]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    output=VGG16_predict(img_path)
    if output>=151 and output<=268 :
        return True
    else:
        return False # true/false
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

```
In [75]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
count_dog=0
for img in human_files_short :
    output= dog_detector(img)
    if output == True :
        count_dog+=1
print("the accuracy of the algorithm in humen files {}".format(count_dog/len(human_files_short)))
```

the accuracy of the algorithm in humen files 0.0

```
In [76]: count_dog=0
for img in dog_files_short :
    output= dog_detector(img)
    if output == True :
        count_dog+=1
print("the accuracy of the algorithm in dog files {}".format((count_dog/len(dog_files_short))))
```

the accuracy of the algorithm in dog files 0.97

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.


```
In [77]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```

In [78]: import os
         from torchvision import datasets
         import torchvision.transforms as transforms

         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         # number of subprocesses to use for data loading
         num_workers = 0
         # how many samples per batch to load
         batch_size = 16

         # convert data to a normalized torch.FloatTensor

         transform_train = transforms.Compose([
             transforms.RandomRotation(45),

             transforms.Resize((512,512)),
             transforms.ToTensor(),
             transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
         ])

         transform_valid = transforms.Compose([
             transforms.Resize((512,512)),
             transforms.ToTensor(),
             transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
         ])

         transform_test = transforms.Compose([
             transforms.Resize((512,512)),
             transforms.ToTensor(),
             transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
         ])

         train_dataset = datasets.ImageFolder(root='/data/dog_images/train',
                                             transform=transform_train)

         test_dataset = datasets.ImageFolder(root='/data/dog_images/test',
                                             transform=transform_test)

         valid_dataset = datasets.ImageFolder(root='/data/dog_images/valid',
                                             transform=transform_valid)

         train_dataloader=torch.utils.data.DataLoader(train_dataset,

```

```

        batch_size=batch_size, shuffle=True,
        num_workers=num_workers)

valid_dataloader=torch.utils.data.DataLoader(valid_dataset,
        batch_size=batch_size,
        num_workers=num_workers)

test_dataloader=torch.utils.data.DataLoader(test_dataset,
        batch_size=batch_size, shuffle=True,
        num_workers=num_workers)

loaders_scratch=[train_dataloader,valid_dataloader,test_dataloader]

# loaders_scratch = {"train" : torch.utils.data.DataLoader(train_dataset,
#
#                                     batch_size=batch_size, shuffle=True,
#                                     num_workers=num_workers)
#
#                                     , "valid" : torch.utils.data.DataLoader(valid_dataset,
#                                     batch_size=batch_size,
#                                     num_workers=num_workers)
#
#                                     , "test" : torch.utils.data.DataLoader(test_dataset,
#                                     batch_size=4, shuffle=True,
#                                     num_workers=num_workers)}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why?

- Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

The first question

I resize the images by stretching to 512. After researching and several attempts, I found that the best size is to be 512 by 512 pixels, but of course I can use a larger scale or less, but if I try to use a greater resolution, this will consume more than the RAM and you will need more time to learn, and if I try to reduce it more, I may reach a fun in which the relationships between Pixels, and you will not be able to identify the type of dog

second question

Yes, I used RandomRotation(Randomize values for the Rotation)

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [79]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        ## input is 512*512*3

        self.pool=nn.MaxPool2d(4,4)

        self.conv1_1=nn.Conv2d(3,16,3,padding=1)
        self.conv1_2=nn.Conv2d(16,16,3,padding=1)
        self.conv1_3=nn.Conv2d(16,16,3,padding=1)

        self.Batch_1=nn.BatchNorm2d(16)

        self.conv2_1=nn.Conv2d(16,32,3,padding=1)
        self.conv2_2=nn.Conv2d(32,32,3,padding=1)
        self.conv2_3=nn.Conv2d(32,32,3,padding=1)

        self.Batch_2=nn.BatchNorm2d(32)

        self.conv3_1=nn.Conv2d(32,64,3,padding=1)
        self.conv3_2=nn.Conv2d(64,64,3,padding=1)
        self.conv3_3=nn.Conv2d(64,64,3,padding=1)

        self.Batch_3=nn.BatchNorm2d(64)

        self.drop=nn.Dropout(p=0.5)
        ## output is 8*8*64
        self.fc1=nn.Linear(8*8*64,2500)
        self.fc2=nn.Linear(2500,1000)
        self.fc3=nn.Linear(1000,133)

        self.Batch_linear=nn.BatchNorm1d(2500)

    def forward(self, x):
        ## Define forward behavior

        x=F.relu(self.conv1_1(x))
        x=F.relu(self.conv1_2(x))
        x=F.relu(self.conv1_3(x))

        x=self.pool(x)
        x=self.Batch_1(x)

```

```

x=F.relu(self.conv2_1(x))
x=F.relu(self.conv2_2(x))
x=F.relu(self.conv2_3(x))

x=self.pool(x)
x=self.Batch_2(x)

x=F.relu(self.conv3_1(x))
x=F.relu(self.conv3_2(x))
x=F.relu(self.conv3_3(x))

x=self.pool(x)
x=self.Batch_3(x)

x=x.view(-1,8*8*64)
x=self.drop(F.relu(self.fc1(x)))
x=self.Batch_linear(x)
x=self.drop(F.relu(self.fc2(x)))
x= self.fc3(x)

return x

##-## You so NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

At first I looked for examples similar to what I wanted to do. At the beginning I applied three main layers, each one containing three convolutional layers the first one extracts raw information such as color and fonts. The second one discovers more complex information, such as the relationship of colors and shapes together, and the last one to identify the shapes of dogs to differentiate them. After each main layer I put 2d BatchNorm to improve the speed of learning and ensure that all data is on the same scale, then I used 3 linear layers for classification and 1d BatchNorm after the first layer to ensure that all data is on the same scale.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [80]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.Adam(model_scratch.parameters(), lr=0.0001)
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model_scratch.pt'.

```
In [81]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0
            valid_loss = 0.0

            #####
            # train the model #
            #####
            model.train()
            for batch_idx, (data, target) in enumerate(loaders[0]):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                optimizer.zero_grad()

                output = model(data)

                loss = criterion(output, target)

                loss.backward()

                optimizer.step()
                ## find the loss and update the model parameters accordingly
                ## record the average training loss, using something like
                ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
                train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            #####
            # validate the model #
            #####
            model.eval()
```

```

for batch_idx, (data, target) in enumerate(loaders[1]):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    output = model(data)
    loss = criterion(output, target)
    ## update the average validation loss
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min :
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.fo
    torch.save(model.state_dict(), save_path)
    valid_loss_min=valid_loss
return model

```

```

In [ ]: # train the model
        model_scratch = train(40, loaders_scratch, model_scratch, optimizer_scratch,
                               criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.831641      Validation Loss: 4.663046
Validation loss decreased (inf --> 4.663046). Saving model ...
Epoch: 2      Training Loss: 4.544766      Validation Loss: 4.406909
Validation loss decreased (4.663046 --> 4.406909). Saving model ...
Epoch: 3      Training Loss: 4.343581      Validation Loss: 4.235242
Validation loss decreased (4.406909 --> 4.235242). Saving model ...
Epoch: 4      Training Loss: 4.197054      Validation Loss: 4.150343
Validation loss decreased (4.235242 --> 4.150343). Saving model ...
Epoch: 5      Training Loss: 4.072188      Validation Loss: 4.041636
Validation loss decreased (4.150343 --> 4.041636). Saving model ...
Epoch: 6      Training Loss: 3.961784      Validation Loss: 3.872406
Validation loss decreased (4.041636 --> 3.872406). Saving model ...
Epoch: 7      Training Loss: 3.845782      Validation Loss: 3.829316
Validation loss decreased (3.872406 --> 3.829316). Saving model ...
Epoch: 8      Training Loss: 3.764273      Validation Loss: 3.748256
Validation loss decreased (3.829316 --> 3.748256). Saving model ...
Epoch: 9      Training Loss: 3.685254      Validation Loss: 3.727593
Validation loss decreased (3.748256 --> 3.727593). Saving model ...

```

```

Epoch: 10      Training Loss: 3.600853      Validation Loss: 3.611072
Validation loss decreased (3.727593 --> 3.611072). Saving model ...
Epoch: 11      Training Loss: 3.507229      Validation Loss: 3.627371
Epoch: 12      Training Loss: 3.426002      Validation Loss: 3.523224
Validation loss decreased (3.611072 --> 3.523224). Saving model ...
Epoch: 13      Training Loss: 3.353586      Validation Loss: 3.453768
Validation loss decreased (3.523224 --> 3.453768). Saving model ...
Epoch: 14      Training Loss: 3.300133      Validation Loss: 3.510188
Epoch: 15      Training Loss: 3.211167      Validation Loss: 3.379157
Validation loss decreased (3.453768 --> 3.379157). Saving model ...
Epoch: 16      Training Loss: 3.152943      Validation Loss: 3.403621
Epoch: 17      Training Loss: 3.091194      Validation Loss: 3.326652
Validation loss decreased (3.379157 --> 3.326652). Saving model ...
Epoch: 18      Training Loss: 3.054357      Validation Loss: 3.398203
Epoch: 19      Training Loss: 2.989785      Validation Loss: 3.275832
Validation loss decreased (3.326652 --> 3.275832). Saving model ...
Epoch: 20      Training Loss: 2.943362      Validation Loss: 3.290088
Epoch: 21      Training Loss: 2.849448      Validation Loss: 3.288027
Epoch: 22      Training Loss: 2.822857      Validation Loss: 3.267239
Validation loss decreased (3.275832 --> 3.267239). Saving model ...
Epoch: 23      Training Loss: 2.768841      Validation Loss: 3.247331
Validation loss decreased (3.267239 --> 3.247331). Saving model ...
Epoch: 24      Training Loss: 2.710224      Validation Loss: 3.262791
Epoch: 25      Training Loss: 2.647130      Validation Loss: 3.242076
Validation loss decreased (3.247331 --> 3.242076). Saving model ...
Epoch: 26      Training Loss: 2.633497      Validation Loss: 3.231065
Validation loss decreased (3.242076 --> 3.231065). Saving model ...
Epoch: 27      Training Loss: 2.580645      Validation Loss: 3.234056
Epoch: 28      Training Loss: 2.535292      Validation Loss: 3.211452
Validation loss decreased (3.231065 --> 3.211452). Saving model ...
Epoch: 29      Training Loss: 2.482041      Validation Loss: 3.176485
Validation loss decreased (3.211452 --> 3.176485). Saving model ...

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [82]: model_scratch.load_state_dict(torch.load('model_scratch.pt'))
```

```

def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.0
    correct = 0.0
    total = 0.0

```



```

model.eval()
for batch_idx, (data, target) in enumerate(loaders[2]):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

# call test function
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.224298

Test Accuracy: 24% (201/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [83]: ## TODO: Specify data loaders
import os
from torchvision import datasets
from PIL import ImageFile

```

```

ImageFile.LOAD_TRUNCATED_IMAGES = True
# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 16
# convert data to a normalized torch.FloatTensor
transform = transforms.Compose([
    transforms.RandomRotation(45),
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
train_dataset = datasets.ImageFolder(root='/data/dog_images/train',
transform=transform)
test_dataset = datasets.ImageFolder(root='/data/dog_images/test',
transform=transform)
valid_dataset = datasets.ImageFolder(root='/data/dog_images/valid',
transform=transform)
data = {"train" : train_dataset, "valid" : valid_dataset, "test" : test_dataset}
loaders_transfer = {0 : torch.utils.data.DataLoader(train_dataset,
batch_size=batch_size, shuffle=True,
num_workers=num_workers)
,1 : torch.utils.data.DataLoader(valid_dataset,
batch_size=batch_size,
num_workers=num_workers)
,2 : torch.utils.data.DataLoader(test_dataset,
batch_size=4, shuffle=True,
num_workers=num_workers)}

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [84]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture

model_transfer = models.vgg16(pretrained=True)

for param in model_transfer.features.parameters():
    param.requires_grad = False

model_transfer.classifier[6]=torch.nn.Linear(4096,133)

```

```
if use_cuda:
    model_transfer = model_transfer.cuda()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: I used the same VGG model with the same weights on the convolutional layers and drop layers but i have used different finel linear layer with 133 output of dog bread because of the small size of my data and the fact that it is close to the data that I was previously trained on VGG model

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [85]: import torch.optim as optim

criterion_transfer = torch.nn.CrossEntropyLoss()
optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [ ]: # train the model
model_transfer=train(50, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1      Training Loss: 3.971319      Validation Loss: 2.657659
Validation loss decreased (inf --> 2.657659). Saving model ...
Epoch: 2      Training Loss: 2.095121      Validation Loss: 1.419118
Validation loss decreased (2.657659 --> 1.419118). Saving model ...
Epoch: 3      Training Loss: 1.418794      Validation Loss: 1.060492
Validation loss decreased (1.419118 --> 1.060492). Saving model ...
Epoch: 4      Training Loss: 1.195815      Validation Loss: 0.943234
Validation loss decreased (1.060492 --> 0.943234). Saving model ...
Epoch: 5      Training Loss: 1.024925      Validation Loss: 0.903657
Validation loss decreased (0.943234 --> 0.903657). Saving model ...
Epoch: 6      Training Loss: 0.932685      Validation Loss: 0.809092
Validation loss decreased (0.903657 --> 0.809092). Saving model ...
Epoch: 7      Training Loss: 0.875713      Validation Loss: 0.797572
Validation loss decreased (0.809092 --> 0.797572). Saving model ...
Epoch: 8      Training Loss: 0.819681      Validation Loss: 0.780811
Validation loss decreased (0.797572 --> 0.780811). Saving model ...
Epoch: 9      Training Loss: 0.746925      Validation Loss: 0.759027
```

```

Validation loss decreased (0.780811 --> 0.759027). Saving model ...
Epoch: 10      Training Loss: 0.737224      Validation Loss: 0.729879
Validation loss decreased (0.759027 --> 0.729879). Saving model ...
Epoch: 11      Training Loss: 0.685718      Validation Loss: 0.748925
Epoch: 12      Training Loss: 0.669677      Validation Loss: 0.725773
Validation loss decreased (0.729879 --> 0.725773). Saving model ...
Epoch: 13      Training Loss: 0.645981      Validation Loss: 0.747540
Epoch: 14      Training Loss: 0.622783      Validation Loss: 0.721736
Validation loss decreased (0.725773 --> 0.721736). Saving model ...
Epoch: 15      Training Loss: 0.593957      Validation Loss: 0.702301
Validation loss decreased (0.721736 --> 0.702301). Saving model ...
Epoch: 16      Training Loss: 0.569698      Validation Loss: 0.728802
Epoch: 17      Training Loss: 0.547321      Validation Loss: 0.690065
Validation loss decreased (0.702301 --> 0.690065). Saving model ...
Epoch: 18      Training Loss: 0.535204      Validation Loss: 0.668215
Validation loss decreased (0.690065 --> 0.668215). Saving model ...
Epoch: 19      Training Loss: 0.522946      Validation Loss: 0.681632
Epoch: 20      Training Loss: 0.510970      Validation Loss: 0.663390
Validation loss decreased (0.668215 --> 0.663390). Saving model ...
Epoch: 21      Training Loss: 0.488621      Validation Loss: 0.672864
Epoch: 22      Training Loss: 0.483402      Validation Loss: 0.655389
Validation loss decreased (0.663390 --> 0.655389). Saving model ...
Epoch: 23      Training Loss: 0.475028      Validation Loss: 0.697595
Epoch: 24      Training Loss: 0.475395      Validation Loss: 0.658401
Epoch: 25      Training Loss: 0.455119      Validation Loss: 0.646250
Validation loss decreased (0.655389 --> 0.646250). Saving model ...
Epoch: 26      Training Loss: 0.442444      Validation Loss: 0.688119
Epoch: 27      Training Loss: 0.428311      Validation Loss: 0.669664
Epoch: 28      Training Loss: 0.401939      Validation Loss: 0.683290
Epoch: 29      Training Loss: 0.398759      Validation Loss: 0.668991
Epoch: 30      Training Loss: 0.402228      Validation Loss: 0.665745
Epoch: 31      Training Loss: 0.381120      Validation Loss: 0.666035

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```

In [86]: model_transfer.load_state_dict(torch.load('model_transfer.pt'))

        test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)

```

Test Loss: 0.670501

Test Accuracy: 79% (667/836)

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [87]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

         # list of class names by index, i.e. a name can be accessed like class_names[0]

data_transfer=data

class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]

def predict_breed_transfer(img_path):
    # load the image and return the predicted breed

    transform = transforms.Compose([
transforms.Resize((224,224)),
transforms.ToTensor(),
transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

    img= Image.open(img_path)
    img_t=transform(img)
    batch_t = torch.unsqueeze(img_t, 0)

    model_transfer = models.vgg16(pretrained=True)

    for param in model_transfer.features.parameters():
        param.requires_grad = False

    model_transfer.classifier[6]=torch.nn.Linear(4096,133)

    model_transfer.load_state_dict(torch.load('model_transfer.pt'))

    output=model_transfer(batch_t)

    _, preds_tensor = torch.max(output, 1)
    preds_tensor=preds_tensor.to(torch.device("cpu"))
    preds = np.squeeze(preds_tensor.numpy())

    for dog_type in range(len(class_names)) :
        if dog_type==preds:
            return class_names[dog_type]

    return None
```



Sample Human Output

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

In [88]: *### TODO: Write your algorithm.*

Feel free to use as many code cells as needed.

```
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    plt.subplots()
    img = Image.open(img_path)
    plt.imshow(img)

    predict=face_detector(img_path)
    if predict:
        human_dog_breed= predict_breed_transfer(img_path)
        plt.title("hello ,human you look like {}".format(human_dog_breed))

    predict=dog_detector(img_path)

    if predict:
        dog_breed= predict_breed_transfer(img_path)
        plt.title("the dog bread is {}".format(dog_breed))
```

```
return "oh It seems that I did not know the content of the image Are you sure it is"
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement)

The results are very good.

1- I think that I can do the exercise of the dog type recognition model more.


2- use more accurate facial recognition .

3- make a transfer that integrates the closest image of the human face of the type of dogs with it and displays it.


```
In [89]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.

         ## suggested code, below
         for file in np.hstack((human_files[:3], dog_files[:3])):
             run_app(file)
```


hello ,humen you look like Dachshund




hello ,humen you look like Dogue de bordeaux




hello ,humen you look like Portuguese water dog



the dog bread is Mastiff



the dog bread is Mastiff



the dog bread is Mastiff

